Flash | Citi Research

09 Sep 2025 18:24:29 ET | 11 pages

NVIDIA Corp (NVDA.O)

AI Infra Summit Takeaways

CITI'S TAKE
Today, NVIDIA's VP of Hyperscale and High-Performance Computing, Ian Buck, reiterated the company's goal to accelerate Gen AI adoption through GPU-powered data centers at AI Infrastructure Conference in Santa Clara, California. While various discussions emerged during the keynote, we highlight the following two: ***1) The new Rubin CPX GPU*** which when combined with Vera and Rubin can deliver up up to $5B in token revenue for every $100M invested (~50x ROI vs GB200 NVL72's ~10x). ***2) NVDA announced its GB300 NVL72 rack-scale system set a new reasoning inference benchmark records vs. GB200 NVL72*** (see body for details). We note by inserting the Rubin CPX into the roadmap NVDA is accelerating its one-year cadence amid the ramping ASIC competition (see our recent note). We also see these announcements as an indication that the age of inference is upon us as reflected by Google's recent comments of 50x+ token increase Y/Y.

Buy	
Price (09 Sep 25 16:00)	US$170.76
Target price	US$200.00
Expected share price return	17.1%
Expected dividend yield	0.0%
Expected total return	17.1%
Market Cap	US$4,149,468M

Atif Malik[AC]

Papa Sylla

NVDA announced the following at AI Infra Summit:

Rubin CPX: A New Class of GPU Designed for Long-Context Inference: NVIDIA Rubin CPX is a new class of GPU purpose-built geared for the highest performance and token revenue for ultra-large context processing. Rubin CPX enables AI systems to handle million-token software coding and generative video with groundbreaking speed and efficiency. Rubin CPX delivers up to 3x faster attention capabilities compared with NVIDIA GB300 NVL72 systems. Per NVDA, the chip works tightly with NVIDIA Vera CPUs and Rubin GPUs inside the new NVIDIA Vera Rubin NVL144 CPX platform, which enables companies to monetize investments at an unprecedented scale with ~50x ROI ($5B in token revenue for every $100M invested).

See Appendix A-1 for Analyst Certification, Important Disclosures and Research Analyst Affiliations

Figure 1. Nvidia's Vera Rubin NVL 144 CPX

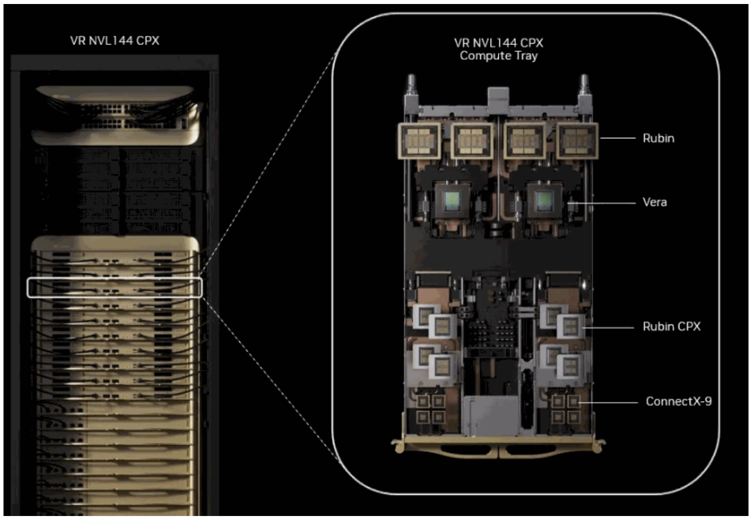

Source: Nvidia

Figure 2. Nvidia's Updated AI Products Cadence

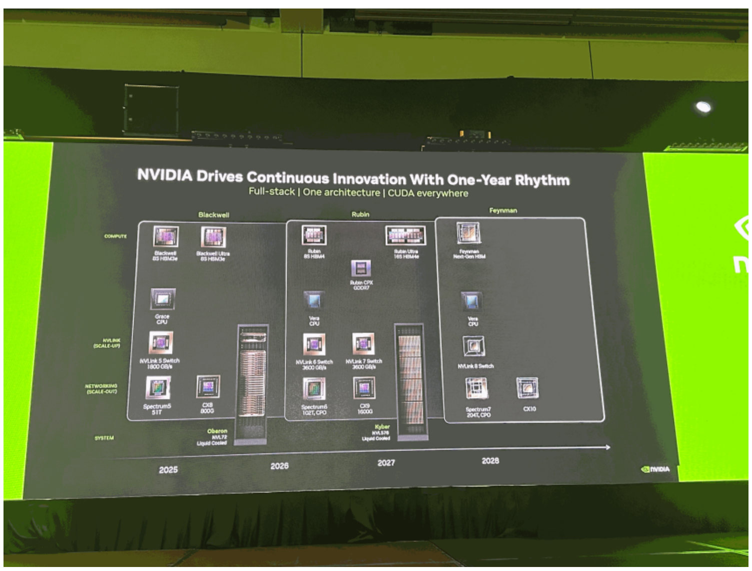

Source: Citi Research, Nvidia's AI Infra Summit Keynote

NVIDIA Blackwell Ultra Sets the Bar in New MLPerf Inference Benchmark: NVDA announced that its NVIDIA GB300 NVL72 rack-scale system set records on the new reasoning inference benchmark in the latest MLPerf inference benchmark. NVDA's newest system can deliver up to 1.4x more DeepSeek-R1 inference throughput compared to GB200 NVL72 systems. The company added that the platform also set performance records on all new data center benchmarks added to the MLPerf Inference v5.1 suite — including DeepSeek-R1, Llama 3.1 405B Interactive, Llama 3.1 8B and Whisper. These system level performance are on top of NVDA's per-GPU records across MLPerf data center benchmark.

NVIDIA Corp

Valuation

Our price target for NVDA of $200 is based on ~30x P/E on C26E EPS. Our 30x P/E multiple is in-line with 3-5 year average.

Risks

Downside risks to the attainment of our target price include: 1) competition on gaming could drive the stock lower if Nvidia loses market share; 2) slower-than-expected adoption of new platforms can drive lower data center and gaming sales; 3) lumpiness in auto and data center markets can add volatility to the stock/multiple; and 4) cryptomining impact on gaming sales.

If you are visually impaired and would like to speak to a Citi representative regarding the details of the graphics in this document, please call USA 1-888-500-5008 (TTY: 711), from outside the US +1-210-677-3788